

- **Vendor: Amazon**
- **Exam Code: SAA-C02**
- **Exam Name: AWS Certified Solutions Architect - Associate**
- **New Questions (Oct/2021)**

[Visit PassLeader and Download Full Version AWS-Associate \(SAA-C02\) Exam Dumps](#)

NEW QUESTION 766

A company hosts more than 300 global websites and applications. The company requires a platform to analyze more than 30 TB of clickstream data each day. What should a solutions architect do to transmit and process the clickstream data?

- A. Design an AWS Data Pipeline to archive the data to an Amazon S3 bucket and run an Amazon EMR cluster with the data to generate analytics.
- B. Create an Auto Scaling group of Amazon EC2 instances to process the data and send it to an Amazon S3 data lake for Amazon Redshift to use for analysis.
- C. Cache the data to Amazon CloudFront. Store the data in an Amazon S3 bucket. When an object is added to the S3 bucket, run an AWS Lambda function to process the data for analysis.
- D. Collect the data from Amazon Kinesis Data Streams. Use Amazon Kinesis Data Firehose to transmit the data to an Amazon S3 data lake. Load the data in Amazon Redshift for analysis.

Answer: D

NEW QUESTION 767

A company hosts a popular website in the AWS Cloud. A solutions architect needs to provide reports about user click behaviour in near-real time as users navigate the website. Which solution will meet this requirement?

- A. Store the clickstream data in Amazon DynamoDB. Deploy an application that runs on AWS Elastic Beanstalk to process and analyze the data.
- B. Push the clickstream data from each session to an Amazon Kinesis data stream. Analyze the data by using Amazon Kinesis Data Analytics.
- C. Store the clickstream data in an Amazon S3 bucket. Order the data by timestamp. Process the data with an AWS Lambda function that is subscribed to object creation events on the S3 bucket.
- D. Forward the clickstream data to Amazon Simple Queue Service (Amazon SQS). Store the data in an Amazon RDS for MySQL DB instance. Deploy Amazon EC2 instances to process and analyze the data.

Answer: B

NEW QUESTION 768

A company plans to host a survey website on AWS. The company anticipates an unpredictable amount of traffic. This traffic results in asynchronous updates to the database. The company wants to ensure that writes to the database hosted on AWS do not get dropped. How should the company write its application to handle these database requests?

- A. Configure the application to publish to an Amazon Simple Notification Service (Amazon SNS) topic. Subscribe the database to the SNS topic.
- B. Configure the application to subscribe to an Amazon Simple Notification Service (Amazon SNS) topic. Publish the database updates to the SNS topic.
- C. Use Amazon Simple Queue Service (Amazon SQS) FIFO queues to queue the database connection until the database has resources to write the data.

D. Use Amazon Simple Queue Service (Amazon SQS) FIFO queues to capture the writes and draining the queue as each write is made to the database.

Answer: D

NEW QUESTION 769

A company stores recordings on a monthly basis. Users access the recorded files randomly within 1 year of recording, but users rarely access the files after 1 year. The company wants to optimize its solution by allowing only files that are newer than 1 year old to be queried and retrieved as quickly as possible. A delay in retrieving older files is acceptable. Which solution meets these requirements MOST cost-effectively?

- A. Store individual files in Amazon S3 Glacier.
Store search metadata in object tags that are created in S3 Glacier.
Query the S3 Glacier tags to retrieve the files from S3 Glacier.
- B. Store individual files in Amazon S3.
Use S3 Lifecycle policies to move the files to S3 Glacier after 1 year.
Query and retrieve the files that are in Amazon S3 by using Amazon Athena.
Query and retrieve the files that are in S3 Glacier by using S3 Glacier Select.
- C. Store individual files in Amazon S3.
Store search metadata for each archive in Amazon S3.
Use S3 Lifecycle policies to move the files to S3 Glacier after 1 year.
Query and retrieve the files by searching for metadata from Amazon S3.
- D. Store individual files in Amazon S3 using S3.
Use Lifecycle policies to move the files to S3 Glacier after 1 year.
Store search metadata in Amazon RDS.
Query the files from Amazon RDS. Retrieve the files from Amazon S3 or S3 Glacier.

Answer: D

NEW QUESTION 770

A company uses a combination of Amazon EC2 instances and AWS Fargate tasks to process daily transactions. The company faces unpredictable and sudden increases in transaction volume. The company needs a solution that will process the transactions immediately. Which solution meets these requirements MOST cost-effectively?

- A. Purchase a Compute Savings Plan.
- B. Purchase an EC2 Instance Savings Plan.
- C. Purchase Reserved Instances for existing EC2 workloads.
- D. Use Spot Instances for existing EC2 workloads.
- E. Use On-Demand Spot capacity for the tasks.

Answer: B

NEW QUESTION 771

A meteorological startup company has a custom web application to sell weather data to its users online. The company uses Amazon DynamoDB to store its data and wants to build a new service that sends an alert to the managers of four internal teams every time a new weather event is recorded. The company does not want the new service to affect the performance of the current application. What should a solutions architect do to meet these requirements with the LEAST amount of operational overhead?

- A. Use DynamoDB transactions to write new event data to the table.
Configure the transactions to notify internal teams.
- B. Have the current application publish a message to four Amazon Simple Notification Service (Amazon SNS) topics.
Have each team subscribe to one topic.
- C. Enable Amazon DynamoDB Streams on the table.
Use triggers to write to a single Amazon Simple Notification Service (Amazon SNS) topic to which the teams can subscribe.
- D. Add a custom attribute to each record to flag new items.
Write a cron job that scans the table every minute for items that are new and notifies an Amazon Simple Queue Service (Amazon SQS) queue to which the teams can subscribe.

Answer: C

NEW QUESTION 772

A company has NFS servers in an on-premises data center that need to periodically back up small amounts of data to Amazon S3. Which solution meets these requirements and is MOST cost-effective?

- A. Set up AWS Glue to copy the data from the on-premises servers to Amazon S3.
- B. Set up an AWS DataSync agent on the on-premises servers, and sync the data to Amazon S3.
- C. Set up an SFTP sync using AWS Transfer for SFTP to sync data from on-premises to Amazon S3.
- D. Set up an AWS Direct Connect connection between the on-premises data center and a VPC, and copy the data to Amazon S3.

Answer: B

NEW QUESTION 773

A company is hosting a web application from an Amazon S3 bucket. The application uses Amazon Cognito as an identity provider to authenticate users and return a JSON Web Token (JWT) that provides access to protected resources that are stored in another S3 bucket. Upon deployment of the application, users report errors and are unable to access the protected content. A solutions architect must resolve this issue by providing proper permissions so that users can access the protected content. Which solution meets these requirements?

- A. Update the Amazon Cognito identity pool to assume the proper IAM role for access to the protected content.
- B. Update the S3 ACL to allow the application to access the protected content.
- C. Redeploy the application to Amazon S3 to prevent eventually consistent reads in the S3 bucket from affecting the ability of users to access the protected content.
- D. Update the Amazon Cognito pool to use custom attribute mappings within the identity pool and grant users the proper permissions to access the protected content.

Answer: B

NEW QUESTION 774

A company needs to save the results from a medical trial to an Amazon S3 repository. The repository must allow a few scientists to add new files and must restrict all other users to read-only access. No users can have the ability to modify or delete any files in the repository. The company must keep every file in the repository for a minimum of 1 year after its creation date. Which solution will meet these requirements?

- A. Use S3 Object Lock in governance mode with a legal hold of 1 year.
 - B. Use S3 Object Lock in compliance mode with a retention period of 365 days.
 - C. Use an IAM role to restrict all users from deleting or changing objects in the S3 bucket.
- Use an S3 bucket policy to only allow the IAM role.
- D. Configure the S3 bucket to invoke an AWS Lambda function every time an object is added.
- Configure the function to track the hash of the saved object so that modified objects can be marked accordingly.

Answer: B

NEW QUESTION 775

A company is migrating its data center and wants to securely transfer 50 TB of data to AWS within 2 weeks. The existing data center has a Site-to-Site VPN connection to AWS that is 90 % utilized. Which AWS service should a solutions architect use to meet these requirements?

- A. AWS DataSync with a VPC endpoint.
- B. AWS Direct Connect.
- C. AWS Snowball Edge Storage Optimized.
- D. AWS Storage Gateway.

Answer: C

NEW QUESTION 776

A company runs an internet-facing web application on AWS. The company uses Amazon Route 53 for DNS management and has a public hosted zone to route traffic from the internet to the application. The company wants to log DNS response codes to help system administrators perform any root cause analysis in the future. Which solution will meet these requirements?

- A. Use Route 53 to configure query logging.
- B. Use AWS CloudTrail to record all Route 53 queries.
- C. Use Amazon CloudWatch to record and process Route 53 metrics.
- D. Use AWS Trusted Advisor to perform on-demand root cause analysis.

Answer: A

NEW QUESTION 777

A company used an AWS Direct Connect connection to copy 1 TB of data from a colocation facility to an Amazon S3 bucket in the us-east-1 Region. The company now wants to copy the data to another S3 bucket in the us-west-2 Region. Which solution will meet this requirement?

- A. Use an AWS Snowball Edge Storage Optimized device to copy the data from the colocation facility to us-west-2.
- B. Use the S3 console to copy the data from the source S3 bucket to the target S3 bucket.
- C. Use S3 Transfer Acceleration and the S3 copy-object command to copy the data from the source S3 bucket to the target S3 bucket.
- D. Add an S3 Cross-Region Replication configuration to copy the data from the source S3 bucket to the target S3 bucket.

Answer: D

NEW QUESTION 778

A solution architect is designing the architecture of a new application being deployed to the AWS Cloud. The application will run on Amazon EC2 On-Demand instances and will automatically scale across multiple Availability Zones. The EC2 instances will scale up and down frequently the day. An Application load balancer (ALB) will handle the load distribution. The architecture needs to support distributed session data management. The company is willing to make charges to code if needed. What should the solutions architect do to ensure that the architecture supports distributed session data management?

- A. Use Amazon ElastiCache to manage and store session data.
- B. Use session affinity (sticky sessions) of the ALB to manage session data.
- C. Use Session Manager from AWS Systems Manager to manage the session.
- D. Use the GetSessionToken API operation in AWS Security Token Service (AWS STS) to manage the session.

Answer: C

NEW QUESTION 779

A company uses an application to present metrics from sporting events to the public. The application must scale quickly during live events and must store these metrics for long-term reporting purposes. The company's architecture includes the following:

- Amazon EC2 instances that run in an Auto Scaling group in private subnets.
- A network Load Balancer That runs in public subnets.
- A MongoDB database cluster that runs across multiple EC2 instances.

A solutions architect must implement a solution that minimizes operational overhead. The solution also must be able to scale automatically. What should the solutions architect set up to meet these requirements?

- A. An Amazon DynamoDB database.
- B. An Amazon RDS for MySQL DB instance.
- C. EC2 instances that run MySQL.
- D. Amazon Redshift.

Answer: A

NEW QUESTION 780

An application hosted on AWS is experiencing performance problems, and the application vendor wants to perform an analysis of the log file to troubleshoot further. The log file is stored on Amazon S3 and is 10GB in size. The application owner will make the log file available to the vendor for a limited time. Which is the MOST secure way to do this?

- A. Enable public read on the S3 object and provide the link to the vendor.
- B. Upload the log to Amazon WorkDocs and share the public link with the vendor.
- C. Generate a presigned URL and have the vendor download the log file before it expires.
- D. Create an IAM user for the vendor to provide access to the S3 bucket and the application. Enforce multifactor authentication.

Answer: C

NEW QUESTION 781

A company collects 10 GB of telemetry data daily from various machines. The company stores the data in an Amazon S3 bucket in a source data account. The company has hired several consulting agencies to use this data for analysis. Each agency needs read access to the data for its analysis. The company must share the data from the source data account by choosing a solution that maximizes security and operational efficiency. Which solution will meet these

requirements?

- A. Configure S3 global tables to replicate data for each agency.
- B. Make the S3 bucket public for a limited time inform only the agencies.
- C. Configure cross-account access for the S3 bucket to the accounts that the agencies own.
- D. Set up an IAM user for each analyst in the source data account. Grant each user access to the S3 bucket.

Answer: A

NEW QUESTION 782

A company is preparing to store confidential data in Amazon S3. For compliance reasons the data must be encrypted at rest Encryption key usage must be logged for auditing purposes. Keys must be rotated every year. Which solution meets these requirements and the MOST operationally efficient?

- A. Server-side encryption with customer-provided keys (SSE-C).
- B. Server-side encryption with Amazon S3 managed keys (SSE-S3).
- C. Server-side encryption with AWS KMS (SSE-KMS) customer master keys (CMKs) with manual rotation.
- D. Server-side encryption with AWS KMS (SSE-KMS) customer master keys (CMKs) with automate rotation.

Answer: D

NEW QUESTION 783

A company has a corporate network on premises and has three VPCs in the AWS Cloud. The company has one VPC each for development, test, and, production. The company wants its system administrators to security gain remote command-line access from the corporate network to Amazon EC2 instances in the VPCs. Which solution meets these requirements MOST cost-effectively?

- A. Set up a VPN connection between the corporate network and each of the three VPCs by using AWS VPN. Use Remote Desktop Protocol (RDP) or SSH over the VPN connection to access the EC2 instances remotely.
- B. Configure the EC2 instances to use an instance profile that trusts AWS Systems Manager. Use Systems Manager Session Manager to gain console access to the EC2 instances.
- C. Create a new VPC Purchase and install a virtual router from AWS Marketplace. Establish a VPN connection from the corporate network to this router. Establish another VPN connection from the 'outer to the other three VPCs. Use Remote Desktop Protocol (RDP) or SSH over the VPN connection to access the EC2 instances remotely.
- D. Create a new VPC Establish a VPN connection to the new VPC. Configure peering connections between the new VPC and the existing VPCs. In the new VPC create an EC2 bastion host to serve as a jump box to EC2 instances in the other VPCs. Use Remote Desktop Protocol (RDP) or SSH over the VPN connection to the bastion host.

Answer: A

NEW QUESTION 784

A solutions architect is designing a VPC with public and private subnets. The VPC and subnets use IPv4 CIDR blocks. There is one public subnet and one private subnet in each of three Availability Zones (AZs) for high availability. An internet gateway is used to provide internet access for the public subnets. The private subnets require access to the internet to allow Amazon EC2 instances to download software updates. What should the solutions architect do to enable internet access for the private subnets?

- A. Create three NAT gateways, one for each public subnet in each AZ. Create a private route table for each AZ that forwards non-VPC traffic to the NAT gateway in its AZ.
- B. Create three NAT instances, one for each private subnet in each AZ. Create a private route table for each AZ that forwards non-VPC traffic to the NAT instance in its AZ.
- C. Create a second internet gateway on one of the private subnets. Update the route table for the private subnets that forward non-VPC traffic to the private Internet gateway.
- D. Create an egress-only internet gateway on one of the public subnets. Update the route table for the private subnets that forward non-VPC traffic to the egress-only internet gateway.

Answer: A

NEW QUESTION 785

A company recently migrated a message processing system to AWS. The system receives messages into an ActiveMQ queue running on an Amazon EC2 instance. Messages are processed by a consumer application running on Amazon EC2. The consumer application processes the messages and writes results to a MySQL database running on Amazon EC2. The company wants this application to be highly available with low operational complexity. Which architecture offers

the HIGHEST availability?

- A. Add a second ActiveMQ server to another Availability Zone.
Add an additional consumer EC2 instance in another Availability Zone.
Replicate the MySQL database to another Availability Zone.
- B. Use Amazon MQ with active/standby brokers configured across two Availability Zones.
Add an additional consumer EC2 instance in another Availability Zone.
Replicate the MySQL database to another Availability Zone.
- C. Use Amazon MQ with active/standby brokers configured across two Availability Zones.
Add an additional consumer EC2 instance in another Availability Zone.
Use Amazon RDS for MySQL with Multi-AZ enabled.
- D. Use Amazon MQ with active/standby brokers configured across two Availability Zones.
Add an Auto Scaling group for the consumer EC2 instances across two Availability Zones.
Use Amazon RDS (or MySQL with Multi-AZ enabled).

Answer: D

NEW QUESTION 786

.....

[Visit PassLeader and Download Full Version AWS-Associate \(SAA-C02\) Exam Dumps](#)