

- **Vendor: Amazon**
- **Exam Code: SAA-C02**
- **Exam Name: AWS Certified Solutions Architect - Associate**
- **New Questions (Sep/2021)**

[Visit PassLeader and Download Full Version AWS-Associate \(SAA-C02\) Exam Dumps](#)

**NEW QUESTION 731**

A developer has a script to generate daily reports that users previously ran manually. The script consistently completes in under 10 minutes. The developer needs to automate this process in a cost-effective manner. Which combination of services should the developer use? (Choose two.)

- A. AWS Lambda.
- B. AWS CloudTrail.
- C. Cron on an Amazon EC2 instance.
- D. Amazon EC2 On-Demand Instance with user data.
- E. Amazon EventBridge (Amazon CloudWatch Events).

Answer: CE

**NEW QUESTION 732**

A solution architect is creating a new Amazon CloudFront distribution for an application. Some of the information submitted by users is sensitive. The application uses HTTPS but needs another layer of security. The sensitive information should be protected throughout the entire application stack and access to the information should be restricted to certain applications. Which action should the solutions architect take?

- A. Configure a CloudFront signed URL.
- B. Configure a CloudFront signed cookie.
- C. Configure a CloudFront field-level encryption profile.
- D. Configure CloudFront and set the Origin Protocol Policy setting to HTTPS Only for the Viewer Protocol Policy.

Answer: C

**NEW QUESTION 733**

A company has an Amazon S3 bucket that contains confidential information in its production AWS account. The company has turned on AWS CloudTrail for the account. The account sends a copy of its logs to Amazon CloudWatch Logs. The company has configured the S3 bucket to log read and write data events. A company auditor discovers that some objects in the S3 bucket have been deleted. A solutions architect must provide the auditor with information about who deleted the objects. What should the solutions architect do to provide this information?

- A. Create a CloudWatch Logs filter to extract the S3 write API calls against the S3 bucket.
- B. Query the CloudTrail logs with Amazon Athena to identify the S3 write API calls against the S3 bucket.
- C. Use AWS Trusted Advisor to perform security checks for S3 write API calls that deleted the content.
- D. Use AWS Config to track configuration changes on the S3 bucket. Use these details to track the S3 write API calls that deleted the content.

Answer: B

**NEW QUESTION 734**

A company has three AWS accounts: Management, Development, and Production. These accounts use AWS services only in the us-east-1 Region. All accounts have a VPC with VPC Flow Logs configured to publish data to an Amazon S3 bucket in each separate account. For compliance reasons the company needs an ongoing method to aggregate all the

VPC flow logs across all accounts into one destination S3 bucket in the Management account. What should a solutions architect do to meet these requirements with the LEAST operational overhead?

- A. Add S3 Same-Region Replication rules in each S3 bucket that stores VPC flow logs to replicate objects to the destination S3 bucket.  
Configure the destination S3 bucket to allow objects to be received from the S3 buckets in other accounts.
- B. Set up an IAM user in the Management account.  
Grant permissions to the IAM user to access the S3 buckets that contain the VPC flow logs.  
Run the aws s3 sync command in the AWS CLI to copy the objects to the destination S3 bucket.
- C. Use an S3 inventory report to specify which objects in the S3 buckets to copy.  
Perform an S3 batch operation to copy the objects into the destination S3 bucket in the Management account with a single request.
- D. Create an AWS Lambda function in the Management account Grant S3 GET permissions on the source S3 buckets.  
Grant S3 PUT permissions on the destination S3 bucket.  
Configure the function to invoke when objects are loaded in the source S3 buckets.

Answer: A

**NEW QUESTION 735**

A company is running a multi-tier web application on AWS. The application runs its database on Amazon Aurora MySQL. The application and database tiers are in the us-east-1 Region. A database administrator who monitors the Aurora DB cluster finds that an intermittent increase in read traffic is creating high CPU utilization on the read replica. The result is increased read latency for the application. The memory and disk utilization of the DB instance are stable throughout the event of increased latency. What should a solutions architect do to improve the read scalability?

- A. Reboot the DB cluster.
- B. Create a cross-Region read replica.
- C. Configure Aurora Auto Scaling for the read replica.
- D. Increase the provisioned read IOPS for the DB instance.

Answer: B

**NEW QUESTION 736**

A developer is creating an AWS Lambda function to perform dynamic updates to a database when an item is added to an Amazon Simple Queue Service (Amazon SQS) queue. A solutions architect must recommend a solution that tracks any usage of database credentials in AWS CloudTrail. The solution also must provide auditing capabilities. Which solution will meet these requirements?

- A. Store the encrypted credentials in a Lambda environment variable.
- B. Create an Amazon DynamoDB table to store the credentials. Encrypt the table.
- C. Store the credentials as a secure string in AWS Systems Manager Parameter Store.
- D. Use an AWS Key Management Service (AWS KMS) key store to store the credentials.

Answer: D

**NEW QUESTION 737**

A company has a service that reads and writes large amounts of data from an Amazon S3 bucket in the same AWS Region. The service is deployed on Amazon EC2 instances within the private subnet of a VPC. The service communicates with Amazon S3 over a NAT gateway in the public subnet. However, the company wants a solution that will reduce the data output costs. Which solution will meet these requirements MOST cost-effectively?

- A. Provision a dedicated EC2 NAT instance in the public subnet.  
Configure the route table for the private subnet to use the elastic network interface of this instance as the destination for all S3 traffic.
- B. Provision a dedicated EC2 NAT instance in the private subnet.  
Configure the route table for the public subnet to use the elastic network interface of this instance as the destination for all S3 traffic.
- C. Provision a VPC gateway endpoint.  
Configure the route table for the private subnet to use the gateway endpoint as the route for all S3 traffic.
- D. Provision a second NAT gateway.  
Configure the route table for the private subnet to use this NAT gateway as the destination for all S3 traffic.

Answer: C

NEW QUESTION 738

A company has an application that uses an Amazon DynamoDB table for low storage. A solutions architect discovers that many requests to the table are not returning the latest data. The company's users have not reported any other issues with database performance. Latency is in an acceptable range. Which design change should the solutions architect recommend?

- A. Add read replicas to the table.
- B. Use a global secondary index (GSI).
- C. Request strongly consistent reads for the table.
- D. Request eventually consistent reads for the table.

Answer: C

NEW QUESTION 739

A company wants to share data that is collected from self-driving cars with the automobile community. The data will be made available from within an Amazon S3 bucket. The company wants to minimize its cost of making this data available to other AWS accounts. What should a solutions architect do to accomplish this goal?

- A. Create an S3 VPC endpoint for the bucket.
- B. Configure the S3 bucket to be a Requester Pays bucket.
- C. Create an Amazon CloudFront distribution in front of the S3 bucket.
- D. Require that the files be accessible only with the use of the BitTorrent protocol.

Answer: A

NEW QUESTION 740

A company recently announced the deployment of its retail website to a global audience. The website runs on multiple Amazon EC2 instances behind an Elastic Load Balancer. The instances run in an Auto Scaling group across multiple Availability Zones. The company wants to provide its customers with different versions of content based on the devices that the customers use to access the website. Which combination of actions should a solutions architect take to meet these requirements? (Choose two.)

- A. Configure Amazon CloudFront to cache multiple versions of the content.
- B. Configure a host header in a Network Load Balancer to forward traffic to different instances.
- C. Configure a Lambda@Edge function to send specific objects to users based on the User-Agent header.
- D. Configure AWS Global Accelerator.

Forward requests to a Network Load Balancer (NLB).

Configure the NLB to set up host-based routing to different EC2 instances.

- E. Configure AWS Global Accelerator.

Forward requests to a Network Load Balancer (NLB).

Configure the NLB to set up path-based routing to different EC2 instances.

Answer: BD

NEW QUESTION 741

A company has developed a new content-sharing application that runs on Amazon Elastic Container Service (Amazon ECS). The application runs on Amazon Linux Docker tasks that use the Amazon EC2 launch type. The application requires a storage solution that has the following characteristics:

- Accessibility (or multiple ECS tasks through bind mounts).
- Resiliency across Availability Zones.
- Burstable throughput of up to 3 Gbps.
- Ability to be scaled up over time.

Which storage solution meets these requirements?

- A. Launch an Amazon FSx for Windows File Server Multi-AZ instance.  
Configure the ECS task definitions to mount the Amazon FSx instance volume at launch.
- B. Launch an Amazon Elastic File System (Amazon EFS) instance.  
Configure the ECS task definitions to mount the EFS instance volume at launch.
- C. Create a Provisioned IOPS SSD (io2) Amazon Elastic Block Store (Amazon EBS) volume with Multi-Attach set to enabled.  
Attach the EBS volume to the ECS EC2 instance. Configure ECS task definitions to mount the EBS instance volume at launch.
- D. Launch an EC2 instance with several Provisioned IOPS SSD (io2) Amazon Elastic Block Store (Amazon EBS) volumes attached in a RAID 0 configuration.  
Configure the EC2 instance as an NFS storage server.

Configure ECS task definitions to mount the volumes at launch.

Answer: B

NEW QUESTION 742

An airline that is based in the United States provides services for routes in North America and Europe. The airline is developing a new read-intensive application that customers can use to find flights on either continent. The application requires strong read consistency and needs scalable database capacity to accommodate changes in user demand. The airline needs the database service to synchronize with the least possible latency between the two continents and to provide a simple failover mechanism to a second AWS Region. Which solution will meet these requirements?

- A. Deploy Microsoft SQL Server on Amazon EC2 instances in a Region in North America. Use SQL Server binary log replication on an EC2 instance in a Region in Europe.
- B. Create an Amazon DynamoDB global table. Add a Region from North America and a Region from Europe to the table. Query data with strongly consistent reads.
- C. Use an Amazon Aurora MySQL global database. Deploy the read-write node in a Region in North America, and deploy read-only endpoints in Regions in North America and Europe. Query data with global read consistency.
- D. Create a subscriber application that uses Amazon Kinesis Data Streams for an Amazon Redshift cluster in a Region in North America. Create a second subscriber application for the Amazon Redshift cluster in a Region in Europe. Process all database modifications through Kinesis Data Streams.

Answer: C

NEW QUESTION 743

A company has a production web application in which users upload documents through a web interface or a mobile app. According to a new regulatory requirement, new documents cannot be modified or deleted after they are stored. What should a solutions architect do to meet this requirement?

- A. Store the uploaded documents in an Amazon S3 bucket with S3 Versioning and S3 Object Lock enabled.
- B. Store the uploaded documents in an Amazon S3 bucket. Configure an S3 Lifecycle policy to archive the documents periodically.
- C. Store the uploaded documents in an Amazon S3 bucket with S3 Versioning enabled. Configure an ACL to restrict all access to read-only.
- D. Store the uploaded documents on an Amazon Elastic File System (Amazon EFS) volume. Access the data by mounting the volume in read-only mode.

Answer: A

NEW QUESTION 744

A company has a Microsoft .NET application that runs on an on-premises Windows Server. The application stores data by using an Oracle Database Standard Edition server. The company is planning a migration to AWS and wants to minimize development changes while moving the application. The AWS application environment should be highly available. Which combination of actions should the company take to meet these requirements? (Choose two.)

- A. Refactor the application as serverless with AWS Lambda functions running .NET Core.
- B. Rehost the application in AWS Elastic Beanstalk with the .NET platform in a Multi-AZ deployment.
- C. Replatform the application to run on Amazon EC2 with the Amazon Linux Amazon Machine Image (AMI).
- D. Use AWS Database Migration Service (AWS DMS) to migrate from the Oracle database to Amazon DynamoDB in a Multi-AZ deployment.
- E. Use AWS Database Migration Service (AWS DMS) to migrate from the Oracle database to Oracle on Amazon RDS in a Multi-AZ deployment.

Answer: AD

NEW QUESTION 745

A company wants to enforce strict security guidelines on accessing AWS Cloud resources as the company migrates production workloads from its data centers. Company management wants all users to receive permissions according to their job roles and functions. Which solution meets these requirements with the LEAST operational overhead?

- A. Create an AWS Single Sign-On deployment. Connect to the on-premises Active Directory to centrally manage users and permissions across the company.

- B. Create an IAM role for each job function.  
Require each employee to call the sts:AssumeRole action in the AWS Management Console to perform their job role.
- C. Create individual IAM user accounts for each employee.  
Create an IAM policy for each job function, and attach the policy to all IAM users based on their job role.
- D. Create individual IAM user accounts for each employee.  
Create IAM policies for each job function.  
Create IAM groups, and attach associated policies to each group.  
Assign the IAM users to a group based on their Job role.

Answer: D

**NEW QUESTION 746**

A company provides machine learning solutions. The company's users need to download large data sets from the company's Amazon S3 bucket. These downloads often take a long time, especially when the users are running many simulations on a subset of those datasets. Users download the datasets to Amazon EC2 instances in the same AWS Region as the S3 bucket. Multiple users typically use the same datasets at the same time. Which solution will reduce the time that is required to access the datasets?

- A. Configure the S3 bucket to use the S3 Standard storage class with S3 Transfer Acceleration activated.
- B. Configure the S3 bucket to use the S3 Intelligent-Tiering storage class with S3 Transfer Acceleration activated.
- C. Create an Amazon Elastic File System (Amazon EFS) network file system.  
Migrate the datasets by using AWS DataSync.
- D. Move the datasets onto a General Purpose SSD (gp3) Amazon Elastic Block Store (Amazon EBS) volume.  
Attach the volume to all the EC2 instances.

Answer: C

**NEW QUESTION 747**

A company needs to retain its AWS CloudTrail logs (or 3 years). The company is enforcing CloudTrail across a set of AWS accounts by using AWS Organizations from the parent account. The CloudTrail target S3 bucket is configured with S3 Versioning enabled. An S3 Lifecycle policy is in place to delete current objects after 3 years. After the fourth year of use of the S3 bucket, the S3 bucket metrics show that the number of objects has continued to rise. However, the number of new CloudTrail logs that are delivered to the S3 bucket has remained consistent. Which solution will delete objects that are older than 3 years in the MOST cost-effective manner?

- A. Configure the organization's centralized CloudTrail trail to expire objects after 3 years.
- B. Configure the S3 Lifecycle policy to delete previous versions as well as current versions.
- C. Create an AWS Lambda function to enumerate and delete objects from Amazon S3 that are older than 3 years.
- D. Configure the parent account as the owner of all objects that are delivered to the S3 bucket.

Answer: B

**NEW QUESTION 748**

A company has a website hosted on AWS. The website is behind an Application Load Balancer (ALB) that is configured to handle HTTP and HTTPS separately. The company wants to forward all requests to the website so that the requests will use HTTPS. What should a solutions architect do to meet this requirement?

- A. Update the ALB's network ACL to accept only HTTPS traffic.
- B. Create a rule that replaces the HTTP in the URL with HTTPS.
- C. Create a listener rule on the ALB to redirect HTTP traffic to HTTPS.
- D. Replace the ALB with a Network Load Balancer configured to use Server Name Indication (SNI).

Answer: C

**NEW QUESTION 749**

A company is deploying an application that processes large quantities of data in batches as needed. The company plans to use Amazon EC2 instances for the workload. The network architecture must support a highly scalable solution and prevent groups of nodes from sharing the same underlying hardware. Which combination of network solutions will meet these requirements? (Choose two.)

- A. Create Capacity Reservations for the EC2 instances to run in a placement group.
- B. Run the EC2 instances in a spread placement group.
- C. Run the EC2 instances in a cluster placement group.
- D. Place the EC2 instances in an EC2 Auto Scaling group.
- E. Run the EC2 instances in a partition placement group.

Answer: BC

NEW QUESTION 750

A company is planning to store sensitive documents in an Amazon S3 bucket. The documents must be encrypted at rest. The company wants to manage the underlying keys that are used for encryption. However, the company does not want to manage the encryption and decryption process. Which solutions will meet these requirements? (Choose two.)

- A. Use server-side encryption with customer-provided encryption keys (SSE-C).
- B. Use client-side encryption with AWS managed keys.
- C. Use server-side encryption with S3 managed encryption keys (SSE-S3).
- D. Use server-side encryption with AWS KMS managed encryption keys (SSE-KMS) with a key policy document that is 40 KB in size.
- E. Use server-side encryption with AWS KMS managed encryption keys (SSE-KMS) that the company uploads to AWS KMS.

Answer: CE

NEW QUESTION 751

.....

[Visit PassLeader and Download Full Version AWS-Associate \(SAA-C02\) Exam Dumps](#)